

Hyunjong Ok

📍 Pohang, Republic of Korea | ✉ hyunjong.ok@postech.ac.kr | 🌐 [Website](#) | 🐙 [GitHub](#) | 🎓 [Google Scholar](#)

Research Interests

I develop evaluation methods for **multimodal and speech-enabled language models**. My work focuses on **diagnostic benchmark design**, auditory and temporal grounding, full-duplex voice-agent evaluation, and **data-efficient evaluation and inference** that preserve model-ranking fidelity at lower cost.

Keywords: multimodal and speech LLMs, benchmark design, full-duplex interaction, efficient evaluation

Research and Industry Experience

AI Research Engineer Intern

Dec. 2025 – May 2026

Upstage

- Developed data-efficient VLM evaluation methods that preserve model rankings while substantially reducing benchmark scale, by analyzing benchmark redundancy, item informativeness, and cross-model generalization across large VLM evaluation suites.

AI Research Engineer Intern

Mar. 2023 – Sep. 2023

Naver Cloud

- Conducted research on knowledge-enhanced multimodal named entity recognition (SCANNER, NAACL 2024) and implemented OCR post-processing components for document-AI pipelines.

Selected Scholarships

Ph.D. Excellence Scholarship (Science and Engineering)

Jun. 2026

Korea Student Aid Foundation

Publications

*Equal contribution.

TempCore: Are Video QA Benchmarks Temporally Grounded? A Frame Selection Sensitivity Analysis and Benchmark

[\[Paper\]](#)

Hyunjong Ok, Jaeho Lee

Conference on Empirical Methods in Natural Language Processing (EMNLP), 2026

Speculative End-Turn Detector for Efficient Speech Chatbot Assistant

[\[Paper\]](#)

Hyunjong Ok, Suho Yoo, Jaeho Lee

The 64th Annual Meeting of the Association for Computational Linguistics (ACL), 2026

Lost in the Prompt Order: Revealing the Limitations of Causal Attention in Language Models

[\[Paper\]](#)

Hyunjong Ok, Jaeho Lee

Findings of the 64th Annual Meeting of the Association for Computational Linguistics (ACL Findings), 2026

AuditoryBench++: Can Language Models Understand Auditory Knowledge without Hearing?

[\[Paper\]](#)

Hyunjong Ok*, Suho Yoo*, Hyeonjun Kim, Jaeho Lee

IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP, Oral), 2026

S2Cap: A Benchmark and a Baseline for Singing Style Captioning

[\[Paper\]](#)

Hyunjong Ok, Jaeho Lee

ACM International Conference on Information and Knowledge Management (CIKM, Resource Paper), 2025

Imagine to Hear: Auditory Knowledge Generation Can Be an Effective Assistant for Language Models

[\[Paper\]](#)

Suho Yoo*, Hyunjong Ok*, Jaeho Lee

Findings of the 63rd Annual Meeting of the Association for Computational Linguistics (ACL Findings), 2025

AudioBERT: Audio Knowledge Augmented Language Model
 Hyunjong Ok*, Suho Yoo*, Jaeho Lee
 IEEE International Conference on Acoustics, Speech, and Signal Processing (**ICASSP**), 2025

[Paper]

Decoding with Limited Teacher Supervision Requires Understanding When to Trust the Teacher
 Hyunjong Ok, Jegwang Ryu, Jaeho Lee
 Conference on Empirical Methods in Natural Language Processing (**EMNLP**), 2024

[Paper]

SCANNER: Knowledge-Enhanced Approach for Robust Multi-modal Named Entity Recognition of Unseen Entities
 Hyunjong Ok, Taeho Kil, Sukmin Seo, Jaeho Lee
 Conference of the North American Chapter of the Association for Computational Linguistics (**NAACL**), 2024

[Paper]

Post-Trained Language Model Adaptive to Extractive Summarization of Long Spoken Documents
 Hyunjong Ok, Seong-Bae Park
 IEEE International Conference on Acoustics, Speech, and Signal Processing (**ICASSP** Grand Challenge), 2023

[Paper]

Understanding the Effects of Distractors on Reasoning Vision-Language Models
 Jiyun Bae, Hyunjong Ok, Sangwoo Mo, Jaeho Lee
 Conference on Empirical Methods in Natural Language Processing (**EMNLP**), 2026

[Paper]

[Preprints]
OmniKVQuant: KV Cache Quantization for Omni-LLMs
 Suho Yoo*, Hyunjong Ok*, Jongmin Choi, Jihoo Jung, Joon Son Chung
 arXiv preprint arXiv:2609.11582, 2026

[Paper]

[Workshops]
FinTree: Financial Dataset Pretrain Transformer Encoder for Relation Extraction
 Hyunjong Ok
 Workshop on Knowledge Discovery from Unstructured Data in Financial Services (**SIGIR Workshop**), 2023

[Paper]

Education

POSTECH <i>Integrated M.S./Ph.D. Program, Graduate School of Artificial Intelligence</i> Advisor: Prof. Jaeho Lee	<i>Feb. 2024 – Present</i> <i>Expected Aug. 2027</i>
Kyung Hee University <i>B.S. in Software Convergence, Data Science</i>	<i>Feb. 2019 – Aug. 2023</i>

Selected Awards

International	
ICASSP 2023 MUG Grand Challenge, Track 2: Extractive Text Summarization <i>Alibaba Cloud</i>	1st Place <i>Feb. 2023</i>
SIGIR 2023 KDF Workshop Shared Task: Relation Extraction <i>J.P. Morgan AI Research</i>	2nd Place <i>Jun. 2023</i>
Kaggle <i>Two Silver medals, one Bronze medal</i>	Competitions Expert <i>2022</i>
Domestic	
The 12th Conference of the Korean Artificial Intelligence Association (CKAIA)	Outstanding Paper <i>Aug. 2025</i>
The 11th Joint Conference of the Korean Artificial Intelligence Association (JKAIA)	Outstanding Paper <i>Nov. 2024</i>
National Defense AI Challenge <i>Korea Ministry of National Defense</i>	4th Place <i>Dec. 2023</i>

Korean AI Challenge: ASR Track <i>Korea Ministry of Science and ICT</i>	4th Place <i>Nov. 2023</i>
Korean AI Challenge: ASR Track (Minister's Award) <i>Korea Ministry of Science and ICT</i>	1st Place <i>Oct. 2022</i>
Legal Judgment Prediction Challenge: Text Classification <i>Law & Company</i>	1st Place <i>Dec. 2022</i>
Genome EC Number Classification AI Challenge <i>Infoboss</i>	2nd Place <i>Dec. 2022</i>
Sentence Type Classification AI Challenge <i>Sungkyunkwan University</i>	4th Place <i>Dec. 2022</i>

Mentoring

Suho Yoo, M.S. Student at KAIST *Aug. 2024 – Sep. 2025*
 Research mentoring leading to publications at ICASSP 2025 and ACL Findings 2025

Academic Service

Conference Reviewer: ACL Rolling Review (ARR), NeurIPS, CIKM

Technical Skills

Programming: Python

Frameworks and Tools: PyTorch, Hugging Face, Docker, Linux, Git, Slurm